

healthcare-data-analysis (1)

December 1, 2025

```
[3]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from datetime import datetime
import warnings
warnings.filterwarnings('ignore')
```

```
[4]: # Set up visualization style
plt.style.use('seaborn-v0_8')
sns.set_palette("husl")

# Load the dataset
df = pd.read_csv('/kaggle/input/healthcare-data/healthcare_dataset.csv')

# Display basic information about the dataset
print("Dataset Shape:", df.shape)
print("\nFirst few rows:")
print(df.head())
```

Dataset Shape: (10000, 15)

First few rows:

	Name	Age	Gender	Blood Type	Medical Condition	\
0	Tiffany Ramirez	81	Female	O-	Diabetes	
1	Ruben Burns	35	Male	O+	Asthma	
2	Chad Byrd	61	Male	B-	Obesity	
3	Antonio Frederick	49	Male	B-	Asthma	
4	Mrs. Brandy Flowers	51	Male	O-	Arthritis	

	Date of Admission	Doctor	Hospital	\
0	2022-11-17	Patrick Parker	Wallace-Hamilton	
1	2023-06-01	Diane Jackson	Burke, Griffin and Cooper	
2	2019-01-09	Paul Baker	Walton LLC	
3	2020-05-02	Brian Chandler	Garcia Ltd	
4	2021-07-09	Dustin Griffin	Jones, Brown and Murray	

Insurance Provider	Billing Amount	Room Number	Admission Type	\
--------------------	----------------	-------------	----------------	---

```

0      Medicare    37490.983364      146     Elective
1  UnitedHealthcare    47304.064845      404     Emergency
2      Medicare    36874.896997      292     Emergency
3      Medicare    23303.322092      480     Urgent
4  UnitedHealthcare    18086.344184      477     Urgent

```

	Discharge Date	Medication	Test Results
0	2022-12-01	Aspirin	Inconclusive
1	2023-06-15	Lipitor	Normal
2	2019-02-08	Lipitor	Normal
3	2020-05-03	Penicillin	Abnormal
4	2021-08-02	Paracetamol	Normal

1 Data exploration

[5]: df.info()

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 15 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Name              10000 non-null   object 
 1   Age               10000 non-null   int64  
 2   Gender            10000 non-null   object 
 3   Blood Type        10000 non-null   object 
 4   Medical Condition 10000 non-null   object 
 5   Date of Admission 10000 non-null   object 
 6   Doctor            10000 non-null   object 
 7   Hospital          10000 non-null   object 
 8   Insurance Provider 10000 non-null   object 
 9   Billing Amount    10000 non-null   float64
 10  Room Number       10000 non-null   int64  
 11  Admission Type   10000 non-null   object 
 12  Discharge Date   10000 non-null   object 
 13  Medication         10000 non-null   object 
 14  Test Results      10000 non-null   object 
dtypes: float64(1), int64(2), object(12)
memory usage: 1.1+ MB

```

[6]: print("\nDataset Info:")
print("\nMissing Values:")
print(df.isnull().sum())

Dataset Info:

Missing Values:

```
Name          0  
Age          0  
Gender        0  
Blood Type    0  
Medical Condition 0  
Date of Admission 0  
Doctor         0  
Hospital        0  
Insurance Provider 0  
Billing Amount   0  
Room Number     0  
Admission Type   0  
Discharge Date    0  
Medication        0  
Test Results      0  
dtype: int64
```

```
[7]: df.head()
```

```
Name  Age  Gender  Blood Type  Medical Condition  \  
0    Tiffany Ramirez  81  Female  O-          Diabetes  
1    Ruben Burns    35  Male    O+          Asthma  
2    Chad Byrd      61  Male    B-          Obesity  
3    Antonio Frederick  49  Male    B-          Asthma  
4    Mrs. Brandy Flowers  51  Male    O-          Arthritis  
  
Date of Admission  Doctor  Hospital  \  
0    2022-11-17  Patrick Parker  Wallace-Hamilton  
1    2023-06-01  Diane Jackson  Burke, Griffin and Cooper  
2    2019-01-09  Paul Baker    Walton LLC  
3    2020-05-02  Brian Chandler  Garcia Ltd  
4    2021-07-09  Dustin Griffin  Jones, Brown and Murray  
  
Insurance Provider  Billing Amount  Room Number  Admission Type  \\  
0    Medicare      37490.983364  146          Elective  
1    UnitedHealthcare  47304.064845  404          Emergency  
2    Medicare      36874.896997  292          Emergency  
3    Medicare      23303.322092  480          Urgent  
4    UnitedHealthcare  18086.344184  477          Urgent  
  
Discharge Date  Medication  Test Results  
0    2022-12-01  Aspirin    Inconclusive  
1    2023-06-15  Lipitor    Normal  
2    2019-02-08  Lipitor    Normal  
3    2020-05-03  Penicillin  Abnormal  
4    2021-08-02  Paracetamol  Normal
```

```
[8]: # Data preprocessing
df['Date of Admission'] = pd.to_datetime(df['Date of Admission'])
df['Discharge Date'] = pd.to_datetime(df['Discharge Date'])
df['Days Hospitalized'] = (df['Discharge Date'] - df['Date of Admission']).dt.
                           days

# Create age groups
def create_age_group(age):
    if age <= 30:
        return '18-30'
    elif age <= 45:
        return '31-45'
    elif age <= 60:
        return '46-60'
    elif age <= 75:
        return '61-75'
    else:
        return '76+'

df['Age Group'] = df['Age'].apply(create_age_group)

print("Dataset loaded and preprocessed successfully!")
print(f"Dataset shape: {df.shape}")
```

Dataset loaded and preprocessed successfully!
Dataset shape: (10000, 17)

2 1. Age Distribution Across Medical Conditions

```
[9]: # Question 1: What is the age distribution of patients across different medical
      ↵conditions?
plt.figure(figsize=(15, 10))

# Box plot
plt.subplot(2, 2, 1)
sns.boxplot(data=df, x='Medical Condition', y='Age')
plt.title('Age Distribution by Medical Condition')
plt.xticks(rotation=45)
plt.ylabel('Age')

# Violin plot
plt.subplot(2, 2, 2)
sns.violinplot(data=df, x='Medical Condition', y='Age')
plt.title('Age Density by Medical Condition')
plt.xticks(rotation=45)
plt.ylabel('Age')
```

```


# Heatmap of age groups vs conditions
plt.subplot(2, 2, 3)
age_condition = pd.crosstab(df['Age Group'], df['Medical Condition'])
sns.heatmap(age_condition, annot=True, fmt='d', cmap='YlOrRd')
plt.title('Age Group vs Medical Condition Distribution')
plt.ylabel('Age Group')

# Statistical summary
plt.subplot(2, 2, 4)
condition_age_stats = df.groupby('Medical Condition')['Age'].agg(['mean', 'median', 'std', 'count']).round(1)
condition_age_stats.plot(kind='bar', ax=plt.gca())
plt.title('Age Statistics by Medical Condition')
plt.xticks(rotation=45)
plt.ylabel('Age')
plt.legend(bbox_to_anchor=(1.05, 1), loc='upper left')

plt.tight_layout()
plt.show()

print("Average Age by Medical Condition:")
print(condition_age_stats)

```



Average Age by Medical Condition:

	mean	median	std	count
Medical Condition				
Arthritis	51.5	52.0	19.4	1650
Asthma	51.4	52.0	19.5	1708
Cancer	51.6	52.0	19.5	1703
Diabetes	51.8	53.0	19.8	1623
Hypertension	50.7	50.0	19.8	1688
Obesity	51.6	52.0	19.5	1628

3 2. Gender Distribution Across Age Groups


```
[10]: # Question 2: How does gender distribution vary across different age groups?
plt.figure(figsize=(15, 6))

# Stacked bar chart
plt.subplot(1, 2, 1)
gender_age = pd.crosstab(df['Age Group'], df['Gender'])
gender_age_pct = gender_age.div(gender_age.sum(axis=1), axis=0) * 100
gender_age_pct.plot(kind='bar', stacked=True, ax=plt.gca())
plt.title('Gender Distribution by Age Group (%)')
plt.xlabel('Age Group')
plt.ylabel('Percentage')
plt.legend(title='Gender')
plt.xticks(rotation=45)

# Heatmap
plt.subplot(1, 2, 2)
sns.heatmap(gender_age, annot=True, fmt='d', cmap='Blues')
plt.title('Gender Distribution by Age Group (Count)')
plt.xlabel('Gender')
plt.ylabel('Age Group')

plt.tight_layout()
plt.show()

print("Gender Distribution by Age Group:")
print(gender_age)
print("\nPercentage Distribution:")
print(gender_age_pct.round(1))
```



Gender Distribution by Age Group:

Gender Female Male

Age Group

Age Group	Female	Male
18-30	966	937
31-45	1118	1093
46-60	1110	1115
61-75	1151	1064
76+	730	716

Percentage Distribution:

Gender Female Male

Age Group

Age Group	Female (%)	Male (%)
18-30	50.8	49.2
31-45	50.6	49.4
46-60	49.9	50.1
61-75	52.0	48.0
76+	50.5	49.5

4 3. Blood Type Distribution by Medical Conditions

```
[11]: # Question 3: What is the blood type distribution among patients with different medical conditions?
      plt.figure(figsize=(16, 10))

      # Stacked bar chart
      plt.subplot(2, 2, 1)
      blood_condition = pd.crosstab(df['Medical Condition'], df['Blood Type'])
      blood_condition.plot(kind='bar', stacked=True, ax=plt.gca())
      plt.title('Blood Type Distribution by Medical Condition')
      plt.xlabel('Medical Condition')
      plt.ylabel('Number of Patients')
```

```

plt.legend(title='Blood Type', bbox_to_anchor=(1.05, 1), loc='upper left')
plt.xticks(rotation=45)

# Heatmap
plt.subplot(2, 2, 2)
blood_condition_pct = blood_condition.div(blood_condition.sum(axis=1), axis=0) * 100
sns.heatmap(blood_condition_pct, annot=True, fmt='.1f', cmap='YlGnBu')
plt.title('Blood Type Distribution by Medical Condition (%)')
plt.xlabel('Blood Type')
plt.ylabel('Medical Condition')


# Blood type overall distribution
plt.subplot(2, 2, 3)
blood_dist = df['Blood Type'].value_counts()
plt.pie(blood_dist.values, labels=blood_dist.index, autopct='%1.1f%%', startangle=90)
plt.title('Overall Blood Type Distribution')

# Most common blood type by condition
plt.subplot(2, 2, 4)
most_common_blood = df.groupby('Medical Condition')['Blood Type'].agg(lambda x: x.mode().iloc[0] if not x.mode().empty else 'Unknown')
most_common_blood.value_counts().plot(kind='bar')
plt.title('Most Common Blood Type for Each Medical Condition')
plt.xlabel('Blood Type')
plt.ylabel('Number of Conditions')

plt.tight_layout()
plt.show()

print("Blood Type Distribution by Medical Condition:")
print(blood_condition)

```



Blood Type Distribution by Medical Condition:

Blood Type	A+	A-	AB+	AB-	B+	B-	O+	O-
Medical Condition								
Arthritis	202	202	219	204	196	186	225	216
Asthma	220	208	199	210	217	218	227	209
Cancer	219	218	208	195	214	221	224	204
Diabetes	197	206	207	214	207	221	174	197
Hypertension	213	212	211	225	211	196	209	211
Obesity	190	192	214	227	199	210	189	207

5 4. Age Group Susceptibility to Medical Conditions

```
[12]: # Question 4: Which age groups are most susceptible to specific medical conditions?
plt.figure(figsize=(15, 8))

# Calculate susceptibility (percentage of age group with condition)
age_condition_pct = pd.crosstab(df['Age Group'], df['Medical Condition'],
                                 normalize='index') * 100


# Heatmap
sns.heatmap(age_condition_pct, annot=True, fmt='.1f', cmap='RdYlBu_r')
plt.title('Medical Condition Prevalence by Age Group (% of Age Group)')
plt.xlabel('Medical Condition')
```

```

plt.ylabel('Age Group')
plt.tight_layout()
plt.show()

# Find most susceptible age group for each condition
print("Most Susceptible Age Group for Each Medical Condition:")
for condition in df['Medical Condition'].unique():
    susceptible_age = age_condition_pct[condition].idxmax()
    percentage = age_condition_pct[condition].max()
    print(f" {condition}: {susceptible_age} ({percentage:.1f}% of this age group)")

```




Most Susceptible Age Group for Each Medical Condition:

- Diabetes: 61-75 (16.7% of this age group)
- Asthma: 61-75 (17.3% of this age group)
- Obesity: 18-30 (16.6% of this age group)
- Arthritis: 46-60 (17.2% of this age group)
- Hypertension: 18-30 (18.1% of this age group)
- Cancer: 76+ (18.0% of this age group)

6 5. Demographics Correlation with Blood Types

```
[13]: # Question 5: How does patient demographics (age, gender) correlate with blood types?  
plt.figure(figsize=(15, 10))  
  
# Age vs Blood Type  
plt.subplot(2, 2, 1)  
sns.boxplot(data=df, x='Blood Type', y='Age')  
plt.title('Age Distribution by Blood Type')  
plt.ylabel('Age')  
  
# Gender vs Blood Type  
plt.subplot(2, 2, 2)  
gender_blood = pd.crosstab(df['Blood Type'], df['Gender'])  
gender_blood.plot(kind='bar', ax=plt.gca())  
plt.title('Gender Distribution by Blood Type')  
plt.xlabel('Blood Type')  
plt.ylabel('Count')  
plt.legend(title='Gender')  
  
# Age Group vs Blood Type  
plt.subplot(2, 2, 3)  
agegroup_blood = pd.crosstab(df['Age Group'], df['Blood Type'])  
sns.heatmap(agegroup_blood, annot=True, fmt='d', cmap='YlOrBr')  
plt.title('Age Group vs Blood Type Distribution')  
plt.xlabel('Blood Type')  
plt.ylabel('Age Group')  
  
# Statistical analysis  
plt.subplot(2, 2, 4)  
blood_age_stats = df.groupby('Blood Type')['Age'].agg(['mean', 'std', 'count']).  
round(1)  
blood_age_stats['mean'].plot(kind='bar', ax=plt.gca())  
plt.title('Average Age by Blood Type')  
plt.xlabel('Blood Type')  
plt.ylabel('Average Age')  
plt.xticks(rotation=45)  
  
plt.tight_layout()  
plt.show()  
  
print("Age Statistics by Blood Type:")  
print(blood_age_stats)  
print("\nGender Distribution by Blood Type:")  
print(gender_blood)
```



Age Statistics by Blood Type:
mean std count

Blood Type	mean	std	count
A+	52.1	19.9	1241
A-	50.8	19.6	1238
AB+	51.4	19.3	1258
AB-	51.8	19.5	1275
B+	52.2	19.6	1244
B-	50.8	19.6	1252
O+	51.4	19.6	1248
O-	51.1	19.5	1244

Gender Distribution by Blood Type:



Gender Female Male

Blood Type	Female	Male
A+	629	612
A-	633	605
AB+	637	621
AB-	671	604
B+	640	604
B-	635	617
O+	641	607
O-	589	655

7 Medical Conditions & Health Analysis

6. Most Common Medical Conditions

```
[14]: # Question 6: What are the most common medical conditions in the dataset?  
plt.figure(figsize=(15, 6))  
  
# Count plot  
plt.subplot(1, 2, 1)  
condition_counts = df['Medical Condition'].value_counts()  
sns.barplot(x=condition_counts.values, y=condition_counts.index,  
             palette='viridis')  
plt.title('Most Common Medical Conditions')  
plt.xlabel('Number of Patients')  
plt.ylabel('Medical Condition')  
  
# Pie chart  
plt.subplot(1, 2, 2)  
plt.pie(condition_counts.values, labels=condition_counts.index, autopct='%.1f%%',  
        startangle=90)  
plt.title('Medical Condition Distribution')  
  
plt.tight_layout()  
plt.show()  
  
print("Medical Condition Frequency:")  
for condition, count in condition_counts.items():  
    percentage = (count / len(df)) * 100  
    print(f" {condition}: {count} patients ({percentage:.1f}%)")
```



Medical Condition Frequency:
Asthma: 1708 patients (17.1%)

Cancer: 1703 patients (17.0%)
 Hypertension: 1688 patients (16.9%)
 Arthritis: 1650 patients (16.5%)
 Obesity: 1628 patients (16.3%)
 Diabetes: 1623 patients (16.2%)


8 7. Medical Conditions Distribution Across Age Groups

```
[15]: # Question 7: How are medical conditions distributed across different age groups?
plt.figure(figsize=(16, 10))

# Stacked bar chart
condition_age_dist = pd.crosstab(df['Age Group'], df['Medical Condition'])
condition_age_dist.plot(kind='bar', stacked=True, figsize=(14, 8))
plt.title('Medical Condition Distribution Across Age Groups')
plt.xlabel('Age Group')
plt.ylabel('Number of Patients')
plt.legend(title='Medical Condition', bbox_to_anchor=(1.05, 1), loc='upper left')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()

print("Medical Condition Distribution by Age Group:")
print(condition_age_dist)
```

<Figure size 1600x1000 with 0 Axes>



Age Group	Medical Condition Distribution by Age Group:					
	Medical Condition	Arthritis	Asthma	Cancer	Diabetes	Hypertension
18-30	300	323	317	304	344	315
31-45	369	380	380	346	389	347
46-60	382	378	392	365	340	368
61-75	370	384	353	370	378	360
76+	229	243	261	238	237	238

9 8. Medical Conditions with Highest Hospitalization Rates

```
[16]: # Question 8: Which medical conditions have the highest hospitalization rates?
condition_hospitalization = df.groupby('Medical Condition').agg({
    'Days Hospitalized': ['mean', 'median', 'count'],
    'Billing Amount': 'mean'
}).round(2)

condition_hospitalization.columns = ['Avg Days Hospitalized', 'Median Days Hospitalized', 'Patient Count', 'Avg Billing Amount']
condition_hospitalization = condition_hospitalization.sort_values('Avg Days Hospitalized', ascending=False)

plt.figure(figsize=(15, 10))

# Average hospitalization days
plt.subplot(2, 2, 1)
sns.barplot(x=condition_hospitalization['Avg Days Hospitalized'], y=condition_hospitalization.index, palette='coolwarm')
plt.title('Average Hospitalization Days by Medical Condition')
plt.xlabel('Average Days Hospitalized')

# Patient count vs hospitalization days
plt.subplot(2, 2, 2)
plt.scatter(condition_hospitalization['Patient Count'], condition_hospitalization['Avg Days Hospitalized'], s=100, alpha=0.7)
for i, condition in enumerate(condition_hospitalization.index):
    plt.annotate(condition, (condition_hospitalization['Patient Count'].iloc[i], condition_hospitalization['Avg Days Hospitalized'].iloc[i]),
                xytext=(5, 5), textcoords='offset points', fontsize=8)
plt.title('Patient Count vs Average Hospitalization Days')
plt.xlabel('Patient Count')
plt.ylabel('Average Days Hospitalized')

# Billing vs hospitalization
```


```

plt.subplot(2, 2, 3)
plt.scatter(condition_hospitalization['Avg Days Hospitalized'], condition_hospitalization['Avg Billing Amount'], s=100, alpha=0.7)
for i, condition in enumerate(condition_hospitalization.index):
    plt.annotate(condition, (condition_hospitalization['Avg Days Hospitalized'].iloc[i], condition_hospitalization['Avg Billing Amount'].iloc[i]),
                 xytext=(5, 5), textcoords='offset points', fontsize=8)
plt.title('Hospitalization Days vs Average Billing Amount')
plt.xlabel('Average Days Hospitalized')
plt.ylabel('Average Billing Amount ($)')

plt.tight_layout()
plt.show()

print("Hospitalization Statistics by Medical Condition:")
print(condition_hospitalization)

```



Hospitalization Statistics by Medical Condition:

Avg Days Hospitalized Median Days Hospitalized \

Medical Condition

Medical Condition	Avg Days Hospitalized	Median Days Hospitalized
Arthritis	15.99	16.0
Diabetes	15.57	16.0
Asthma	15.48	16.0

Cancer	15.48	15.0
Hypertension	15.43	15.0
Obesity	15.42	15.0

	Patient Count	Avg Billing Amount
Medical Condition		
Arthritis	1650	25187.63
Diabetes	1623	26060.12
Asthma	1708	25416.87
Cancer	1703	25539.10
Hypertension	1688	25198.03
Obesity	1628	25720.84

10 Hospital & Healthcare Provider Analysis

11 11. Hospitals with Highest Patient Volume

```
[17]: # Question 11: Which hospitals have the highest patient volume?  
hospital_stats = df.groupby('Hospital').agg({  
    'Name': 'count',  
    'Billing Amount': ['sum', 'mean'],  
    'Days Hospitalized': 'mean'  
) .round(2)  
  
hospital_stats.columns = ['Patient Count', 'Total Billing', 'Avg Billing', 'Avg  
    ↪Hospitalization Days']  
hospital_stats = hospital_stats.sort_values('Patient Count', ascending=False)  
  
plt.figure(figsize=(15, 10))  
  
# Patient count  
plt.subplot(2, 2, 1)  
sns.barplot(x=hospital_stats['Patient Count'].head(10), y=hospital_stats.  
    ↪head(10).index, palette='viridis')  
plt.title('Top 10 Hospitals by Patient Volume')  
plt.xlabel('Number of Patients')  
  
# Total billing  
plt.subplot(2, 2, 2)  
sns.barplot(x=hospital_stats['Total Billing'].head(10), y=hospital_stats.  
    ↪head(10).index, palette='plasma')  
plt.title('Top 10 Hospitals by Total Billing Amount')  
plt.xlabel('Total Billing Amount ($)')  
  
# Average billing  
plt.subplot(2, 2, 3)
```


```

sns.barplot(x=hospital_stats['Avg Billing'].head(10), y=hospital_stats.head(10).index, palette='coolwarm')
plt.title('Top 10 Hospitals by Average Billing Amount')
plt.xlabel('Average Billing Amount ($)')

plt.tight_layout()
plt.show()

print("Top 10 Hospitals by Patient Volume:")
print(hospital_stats[['Patient Count', 'Total Billing', 'Avg Billing']].head(10))

```




Top 10 Hospitals by Patient Volume:

Hospital	Patient Count	Total Billing	Avg Billing
Smith PLC	19	432283.55	22751.77
Smith and Sons	17	477638.88	28096.40
Smith Ltd	14	428163.07	30583.08
Smith Inc	14	351463.89	25104.56
Johnson PLC	13	269777.54	20752.12
Smith Group	12	247635.08	20636.26
Williams Inc	12	245434.30	20452.86
Williams LLC	12	327522.47	27293.54
Thomas Group	11	327045.37	29731.40

Johnson Ltd 11 305971.19 27815.56

12 12. Patient Distribution Across Hospitals

```
[18]: # Question 12: What is the patient distribution across different hospitals?  
plt.figure(figsize=(15, 6))  
  
# Pie chart for top 10 hospitals  
plt.subplot(1, 2, 1)  
top_hospitals = hospital_stats.head(10)  
plt.pie(top_hospitals['Patient Count'], labels=top_hospitals.index, autopct='%.1f%%', startangle=90)  
plt.title('Patient Distribution - Top 10 Hospitals')  
  
# All hospitals patient count  
plt.subplot(1, 2, 2)  
sns.barplot(x=hospital_stats['Patient Count'], y=hospital_stats.index, palette='viridis')  
plt.title('Patient Distribution Across All Hospitals')  
plt.xlabel('Number of Patients')  
  
plt.tight_layout()  
plt.show()  
  
print("Patient Distribution Summary:")  
print(f"Total hospitals: {len(hospital_stats)}")  
print(f"Average patients per hospital: {hospital_stats['Patient Count'].mean():.1f}")  
print(f"Hospital with most patients: {hospital_stats['Patient Count'].idxmax()}\n\t({hospital_stats['Patient Count'].max()} patients)")  
print(f"Hospital with fewest patients: {hospital_stats['Patient Count'].idxmin()}\n\t({hospital_stats['Patient Count'].min()} patients)")
```



Patient Distribution Summary:
 Total hospitals: 8639
 Average patients per hospital: 1.2
 Hospital with most patients: Smith PLC (19 patients)
 Hospital with fewest patients: Payne-Bailey (1 patients)

13 13. Doctor-to-Patient Ratio

```
[19]: # Question 13: What is the doctor-to-patient ratio across different hospitals?
doctor_stats = df.groupby(['Hospital', 'Doctor']).size().
    ↪reset_index(name='Patient Count')
hospital_doctor_stats = doctor_stats.groupby('Hospital').agg({
    'Doctor': 'count',
    'Patient Count': 'sum'
}).rename(columns={'Doctor': 'Doctor Count'})

hospital_doctor_stats['Patients per Doctor'] = (hospital_doctor_stats['Patient Count'] /
                                                hospital_doctor_stats['Doctor Count']).round(1)
hospital_doctor_stats = hospital_doctor_stats.sort_values('Patients per Doctor', ascending=False)

plt.figure(figsize=(15, 8))

# Patients per doctor
plt.subplot(1, 2, 1)
sns.barplot(x=hospital_doctor_stats['Patients per Doctor'].head(15),
            y=hospital_doctor_stats.head(15).index, palette='RdYlBu_r')
plt.title('Patients per Doctor Ratio by Hospital (Top 15)')
plt.xlabel('Patients per Doctor')


# Doctor count vs patient count
plt.subplot(1, 2, 2)
plt.scatter(hospital_doctor_stats['Doctor Count'],
            ↪hospital_doctor_stats['Patient Count'], alpha=0.7, s=100)
for i, hospital in enumerate(hospital_doctor_stats.index):
    plt.annotate(hospital, (hospital_doctor_stats['Doctor Count'].iloc[i],
                           ↪hospital_doctor_stats['Patient Count'].iloc[i]),
                 xytext=(5, 5), textcoords='offset points', fontsize=8, alpha=0.
    ↪7)
plt.title('Doctor Count vs Patient Count by Hospital')
plt.xlabel('Number of Doctors')
plt.ylabel('Number of Patients')
```

```

plt.tight_layout()
plt.show()

print("Doctor-to-Patient Ratio by Hospital:")
print(hospital_doctor_stats[['Doctor Count', 'Patient Count', 'Patients per Doctor']].sort_values('Patients per Doctor', ascending=False).head(10))

```



Doctor-to-Patient Ratio by Hospital:

Hospital	Doctor Count	Patient Count	Patients per Doctor
Abbott Inc	1	1	1.0
Acosta-Holmes	1	1	1.0
Abbott, Curry and Moore	1	1	1.0
Abbott-Jordan	1	1	1.0
Abbott-Phillips	1	1	1.0
Abbott-Shea	1	1	1.0
Acevedo LLC	1	1	1.0
Acevedo and Sons	1	1	1.0
Acevedo, Rojas and Smith	1	1	1.0
Acosta PLC	1	1	1.0

14 14. Average Billing by Medical Condition

```

[20]: # Question 14: What is the average billing amount for different medical conditions?

condition_billing = df.groupby('Medical Condition')['Billing Amount'].agg(['mean', 'median', 'std', 'count']).round(2)

```

```

condition_billing = condition_billing.sort_values('mean', ascending=False)

plt.figure(figsize=(15, 10))

# Average billing by condition
plt.subplot(2, 2, 1)
sns.barplot(x=condition_billing['mean'], y=condition_billing.index, □
             palette='coolwarm')
plt.title('Average Billing Amount by Medical Condition')
plt.xlabel('Average Billing Amount ($')


# Box plot of billing by condition
plt.subplot(2, 2, 2)
sns.boxplot(data=df, x='Billing Amount', y='Medical Condition', □
             palette='viridis')
plt.title('Billing Amount Distribution by Medical Condition')
plt.xlabel('Billing Amount ($')

# Billing statistics
plt.subplot(2, 2, 3)
condition_billing[['mean', 'median']].plot(kind='bar', ax=plt.gca())
plt.title('Mean vs Median Billing by Medical Condition')
plt.xlabel('Medical Condition')
plt.ylabel('Billing Amount ($)')
plt.xticks(rotation=45)
plt.legend()

plt.tight_layout()
plt.show()

print("Billing Statistics by Medical Condition:")
print(condition_billing)

```



Billing Statistics by Medical Condition:

Medical Condition	mean	median	std	count
Diabetes	26060.12	26162.20	14013.92	1623
Obesity	25720.84	25365.02	14040.79	1628
Cancer	25539.10	25610.64	14081.99	1703
Asthma	25416.87	25073.45	14346.78	1708
Hypertension	25198.03	24920.46	14137.14	1688
Arthritis	25187.63	24739.17	13765.17	1650

15 15. Billing Amount vs Length of Stay

```
[21]: # Question 16: How does billing amount correlate with length of stay?
plt.figure(figsize=(15, 10))

# Scatter plot
plt.subplot(2, 2, 1)
sns.scatterplot(data=df, x='Days Hospitalized', y='Billing Amount', alpha=0.6)
plt.title('Billing Amount vs Days Hospitalized')
plt.xlabel('Days Hospitalized')
plt.ylabel('Billing Amount ($)')

# By medical condition
```

```

plt.subplot(2, 2, 2)
sns.scatterplot(data=df, x='Days Hospitalized', y='Billing Amount', ↴
    hue='Medical Condition', alpha=0.6)
plt.title('Billing Amount vs Days Hospitalized (by Condition)')
plt.xlabel('Days Hospitalized')
plt.ylabel('Billing Amount ($)')
plt.legend(bbox_to_anchor=(1.05, 1), loc='upper left')

# Correlation heatmap
plt.subplot(2, 2, 3)
correlation_matrix = df[['Billing Amount', 'Days Hospitalized', 'Age', 'Room Number']].corr()
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', center=0, ↴
    square=True)
plt.title('Correlation Matrix: Billing vs Other Variables')


# Billing per day
df['Billing per Day'] = df['Billing Amount'] / df['Days Hospitalized']
billing_per_day_by_condition = df.groupby('Medical Condition')['Billing per Day'].mean().sort_values(ascending=False)

plt.subplot(2, 2, 4)
sns.barplot(x=billing_per_day_by_condition.head(10).values, ↴
    y=billing_per_day_by_condition.head(10).index, palette='RdYlBu')
plt.title('Top 10 Conditions by Billing per Hospitalization Day')
plt.xlabel('Average Billing per Day ($)')

plt.tight_layout()
plt.show()

# Correlation statistics
correlation = df['Billing Amount'].corr(df['Days Hospitalized'])
print(f"Correlation between Billing Amount and Days Hospitalized: {correlation:.3f}")
print(f"Average billing per day: ${df['Billing per Day'].mean():.2f}")

```



Correlation between Billing Amount and Days Hospitalized: -0.014
 Average billing per day: \$3321.20

16 16. Insurance Providers Covering Highest Billing Amounts

```
[22]: # Question 17: Which insurance providers cover the highest billing amounts?
insurance_billing = df.groupby('Insurance Provider').agg({
    'Billing Amount': ['sum', 'mean', 'count'],
    'Days Hospitalized': 'mean'
}).round(2)

insurance_billing.columns = ['Total Billing', 'Average Billing', 'Patient Count', 'Avg Hospitalization Days']
insurance_billing = insurance_billing.sort_values('Total Billing', ascending=False)

plt.figure(figsize=(15, 10))

# Total billing by insurance
plt.subplot(2, 2, 1)
sns.barplot(x=insurance_billing['Total Billing'], y=insurance_billing.index, palette='plasma')
plt.title('Total Billing Amount by Insurance Provider')
```

```

plt.xlabel('Total Billing Amount ($)')


# Average billing by insurance
plt.subplot(2, 2, 2)
sns.barplot(x=insurance_billing['Average Billing'], y=insurance_billing.index,
             palette='coolwarm')
plt.title('Average Billing Amount by Insurance Provider')
plt.xlabel('Average Billing Amount ($)')

# Patient count vs total billing
plt.subplot(2, 2, 3)
plt.scatter(insurance_billing['Patient Count'], insurance_billing['Total Billing'],
            s=100, alpha=0.7)
for i, provider in enumerate(insurance_billing.index):
    plt.annotate(provider, (insurance_billing['Patient Count'].iloc[i],
                           insurance_billing['Total Billing'].iloc[i]),
                 xytext=(5, 5), textcoords='offset points', fontsize=8)
plt.title('Patient Count vs Total Billing by Insurance Provider')
plt.xlabel('Patient Count')
plt.ylabel('Total Billing ($)')

plt.tight_layout()
plt.show()

print("Insurance Providers by Billing Amount:")
print(insurance_billing)

```



Insurance Providers by Billing Amount:

Insurance Provider	Total Billing	Average Billing	Patient Count
Cigna	52340171.59	25656.95	2040
Aetna	52321794.76	25837.92	2025
Blue Cross	52125858.90	25652.49	2032
UnitedHealthcare	50250467.70	25404.69	1978
Medicare	48129774.83	25002.48	1925

Avg Hospitalization Days

Insurance Provider	Avg Hospitalization Days
Cigna	15.48
Aetna	15.62
Blue Cross	15.57
UnitedHealthcare	15.34
Medicare	15.80

17 17. Average Billing by Admission Type

```
[23]: # Question 18: What is the average billing amount by admission type?
admission_billing = df.groupby('Admission Type').agg({
    'Billing Amount': ['mean', 'median', 'sum', 'count'],
    'Days Hospitalized': 'mean'
```

```

}).round(2)

admission_billing.columns = ['Avg Billing', 'Median Billing', 'Total Billing', 'Patient Count', 'Avg Hospitalization Days']

plt.figure(figsize=(15, 8))


# Average billing by admission type
plt.subplot(1, 2, 1)
sns.barplot(x=admission_billing['Avg Billing'], y=admission_billing.index, palette='viridis')
plt.title('Average Billing Amount by Admission Type')
plt.xlabel('Average Billing Amount ($)')

# Box plot
plt.subplot(1, 2, 2)
sns.boxplot(data=df, x='Admission Type', y='Billing Amount', palette='Set2')
plt.title('Billing Amount Distribution by Admission Type')
plt.xlabel('Admission Type')
plt.ylabel('Billing Amount ($)')
plt.xticks(rotation=45)

plt.tight_layout()
plt.show()

print("Billing Statistics by Admission Type:")
print(admission_billing)

```



Billing Statistics by Admission Type:				
	Avg Billing	Median Billing	Total Billing	Patient Count \
Admission Type				
Elective	25891.83	26120.77	83941321.51	3242
Emergency	24708.51	23900.17	83193559.68	3367
Urgent	25960.83	25850.43	88033186.59	3391

Avg Hospitalization Days	
Admission Type	
Elective	15.60
Emergency	15.61
Urgent	15.48

18 Admission & Stay Analysis


```
[24]: # Question 20: What are the most common admission types for different medical conditions?
plt.figure(figsize=(16, 10))

# Admission type by condition
admission_condition = pd.crosstab(df['Medical Condition'], df['Admission Type'])

# Stacked bar chart
admission_condition.plot(kind='bar', stacked=True, figsize=(14, 8))
plt.title('Admission Type Distribution by Medical Condition')
plt.xlabel('Medical Condition')
plt.ylabel('Number of Patients')
plt.legend(title='Admission Type', bbox_to_anchor=(1.05, 1), loc='upper left')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()

print("Most Common Admission Type for Each Medical Condition:")
for condition in df['Medical Condition'].unique():
    common_admission = df[df['Medical Condition'] == condition]['Admission Type'].mode().iloc[0]
    count = (df[df['Medical Condition'] == condition]['Admission Type'] == common_admission).sum()
    total = len(df[df['Medical Condition'] == condition])
    percentage = (count / total) * 100
    print(f" {condition}: {common_admission} ({percentage:.1f}%)")
```

<Figure size 1600x1000 with 0 Axes>



Most Common Admission Type for Each Medical Condition:

- Diabetes: Emergency (34.3%)
- Asthma: Urgent (34.1%)
- Obesity: Emergency (35.0%)
- Arthritis: Elective (34.5%)
- Hypertension: Urgent (35.2%)
- Cancer: Emergency (33.9%)

19 Medication & Treatment Analysis

```
[25]: # Question 27: What are the most commonly prescribed medications?
plt.figure(figsize=(15, 8))

# Medication distribution
medication_counts = df['Medication'].value_counts()

# Bar chart
plt.subplot(1, 2, 1)
sns.barplot(x=medication_counts.values, y=medication_counts.index,
            palette='viridis')
plt.title('Most Commonly Prescribed Medications')
plt.xlabel('Number of Prescriptions')

# Pie chart
plt.subplot(1, 2, 2)
```


```

plt.pie(medication_counts.values, labels=medication_counts.index, autopct='%.1f%%', startangle=90)
plt.title('Medication Distribution')

plt.tight_layout()
plt.show()

print("Medication Prescription Frequency:")
for medication, count in medication_counts.items():
    percentage = (count / len(df)) * 100
    print(f" {medication}: {count} prescriptions ({percentage:.1f}%)")

```



Medication Prescription Frequency:

Penicillin: 2079 prescriptions (20.8%)
 Lipitor: 2015 prescriptions (20.2%)
 Ibuprofen: 1976 prescriptions (19.8%)
 Aspirin: 1968 prescriptions (19.7%)
 Paracetamol: 1962 prescriptions (19.6%)

20 29. Relationship Between Medications and Test Results

```

[ ]: # Question 29: Is there a relationship between medications and test results?
plt.figure(figsize=(15, 10))

# Test results by medication
medication_test = pd.crosstab(df['Medication'], df['Test Results'], normalize='index') * 100

```

```

# Stacked bar chart
medication_test.plot(kind='bar', stacked=True, figsize=(14, 8))
plt.title('Test Results Distribution by Medication (%)')
plt.xlabel('Medication')
plt.ylabel('Percentage')
plt.legend(title='Test Results', bbox_to_anchor=(1.05, 1), loc='upper left')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()

# Abnormal test rate by medication
abnormal_by_medication = (df[df['Test Results'] == 'Abnormal'] .
    ↪groupby('Medication').size() /
    df.groupby('Medication').size() * 100) .
    ↪sort_values(ascending=False)

print("Medications with Highest Abnormal Test Results Rate:")
for medication, rate in abnormal_by_medication.head().items():
    total_patients = len(df[df['Medication'] == medication])
    print(f" {medication}: {rate:.1f}% abnormal results ({total_patients} patients)")

```

21 COMPREHENSIVE SUMMARY

```

[27]: print("\n" + "=" * 80)
print("COMPREHENSIVE ANALYSIS SUMMARY")
print("=" * 80)

# Key findings summary
print("\nKEY FINDINGS SUMMARY:")
print(f"1. Dataset Overview: {len(df)} patient records from {df['Date of Admission'].min().strftime('%Y-%m-%d')} to {df['Date of Admission'].max().strftime('%Y-%m-%d')}")
print(f"2. Total Billing: ${df['Billing Amount'].sum():,.2f}")
print(f"3. Most Common Condition: {df['Medical Condition'].value_counts().index[0]} ({df['Medical Condition'].value_counts().iloc[0]} patients)")
print(f"4. Average Hospital Stay: {df['Days Hospitalized'].mean():.1f} days")
print(f"5. Gender Distribution: {df['Gender'].value_counts().to_dict()}")
print(f"6. Top Insurance Provider: {df['Insurance Provider'].value_counts().index[0]}")
print(f"7. Most Prescribed Medication: {df['Medication'].value_counts().index[0]}")
print(f"8. Test Results: {df['Test Results'].value_counts().to_dict()}")

# Financial insights

```

```

print("\nFINANCIAL INSIGHTS:")
print(f"• Highest average billing condition: {df.groupby('Medical Condition')['Billing Amount'].mean().idxmax()} ({df.groupby('Medical Condition')['Billing Amount'].mean().max():.2f})")
print(f"• Most profitable condition: {df.groupby('Medical Condition')['Billing Amount'].sum().idxmax()} ({df.groupby('Medical Condition')['Billing Amount'].sum().max():.2f})")
print(f"• Hospital with highest revenue: {df.groupby('Hospital')['Billing Amount'].sum().idxmax()} ({df.groupby('Hospital')['Billing Amount'].sum().max():.2f})")

# Operational insights
print("\nOPERATIONAL INSIGHTS:")
print(f"• Longest average stay condition: {df.groupby('Medical Condition')['Days Hospitalized'].mean().idxmax()} ({df.groupby('Medical Condition')['Days Hospitalized'].mean().max():.1f} days)")
print(f"• Most common admission type: {df['Admission Type'].value_counts().index[0]} ({df['Admission Type'].value_counts().iloc[0]} cases)")
print(f"• Busiest hospital: {df['Hospital'].value_counts().idxmax()}")
print(f"• {df['Hospital'].value_counts().max()} patients)")

# Clinical insights
print("\nCLINICAL INSIGHTS:")
print(f"• Most common blood type: {df['Blood Type'].value_counts().index[0]}")
print(f"• Average patient age: {df['Age'].mean():.1f} years")

print("\n" + "=" * 80)
print("ANALYSIS COMPLETED SUCCESSFULLY!")
print("=" * 80)

```

=====

COMPREHENSIVE ANALYSIS SUMMARY

=====

KEY FINDINGS SUMMARY:

1. Dataset Overview: 10,000 patient records from 2018-10-30 to 2023-10-30
2. Total Billing: \$255,168,067.78
3. Most Common Condition: Asthma (1708 patients)
4. Average Hospital Stay: 15.6 days
5. Gender Distribution: {'Female': 5075, 'Male': 4925}
6. Top Insurance Provider: Cigna
7. Most Prescribed Medication: Penicillin
8. Test Results: {'Abnormal': 3456, 'Inconclusive': 3277, 'Normal': 3267}

FINANCIAL INSIGHTS:

- Highest average billing condition: Diabetes (\$26060.12)
- Most profitable condition: Cancer (\$43,493,080.71)
- Hospital with highest revenue: Smith and Sons (\$477,638.88)

OPERATIONAL INSIGHTS:

- Longest average stay condition: Arthritis (16.0 days)
- Most common admission type: Urgent (3391 cases)
- Busiest hospital: Smith PLC (19 patients)

CLINICAL INSIGHTS:

- Most common blood type: AB- (1275 patients)
- Average patient age: 51.5 years

=====

ANALYSIS COMPLETED SUCCESSFULLY!

=====